



Operating System-Level Virtualization

Seyyed Ahmad Javadi

sajavadi@aut.ac.ir

Spring 2024



A quick review

Virtualization approach	Method	Performance (boot time, runtime, ..)	Isolation
Hardware-level	Binary translation		
	Paravirtualization		
	Hardware-assisted		
OS-level	Containers		

Operating System-level Virtualization

- Creating **different** and **separated** execution environments for applications that are managed concurrently.
- OS kernel allows for multiple isolated user space instances.

	Is there hypervisor?	How many OSs are involved?
Hardware-level	YES	Multiple OSs
OS-level	NO	Single OS

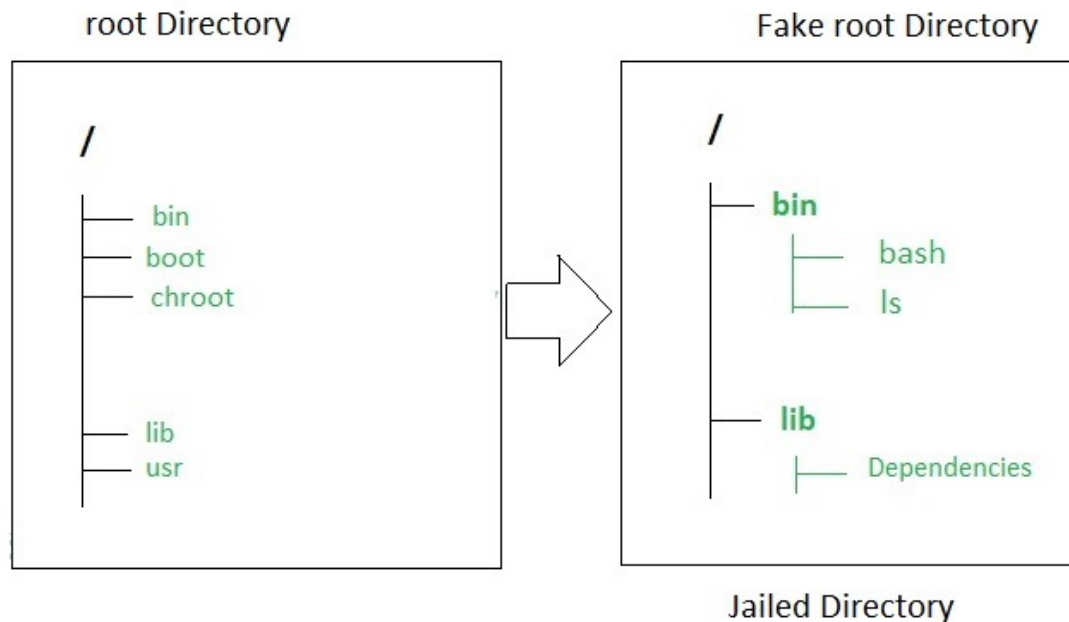
Operating System-level Virtualization (cont.)

- The kernel is also responsible *for*
 - *Sharing the system resources among instances*
 - *Limiting the impact of instances on each other.*
- A *user space instance* contains a proper view of
 - the file system, which is completely isolated
 - separate IP addresses
 - software configurations
 - access to devices.

Operating System-level Virtualization (cont.)

➤ An evolution of the ***chroot mechanism in Unix systems.***

- Changes the file system root directory for a process and its children.
- The process and its children ***cannot have access to other portions of the file system than those accessible under the new root directory.***



<https://www.geeksforgeeks.org/chroot-command-in-linux-with-examples/>

Operating System-level Virtualization (cont.)

- Unix systems ***expose devices as parts of the file system***
 - Using ***chroot*** it is possible to completely ***isolate a set of processes***
- Following the same principle, operating system-level virtualization aims to ***provide separated and multiple execution containers*** for running applications.



Operating System-level Virtualization (cont.)

➤ An ***efficient solution*** for server consolidation scenarios in which ***multiple application servers share the same technology***:

- Operating system
- Application server framework
- Other components.

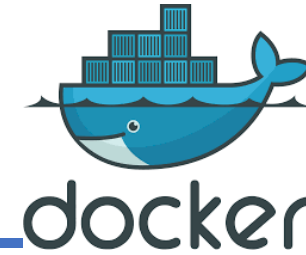
Operating System-level Virtualization (cont.)

- When different **application servers** *are aggregated* into *one physical server*, each server is run in a *different user space*, completely **isolated** from the others.
- Examples of operating system-level virtualizations are:
 - FreeBSD Jails
 - IBM Logical Partition (LPAR)
 - SolarisZones
 - Containers and Docker.

Containers

- OS-level virtualization also called **containerization**.
- A ***container is an isolated virtual env.*** which can run an application.
- Several containers can be created on each operating system, to each of which a ***subset of the computer's resources*** is allocated.
- Programs running inside a container **can only see the container's contents** and devices assigned to the container.

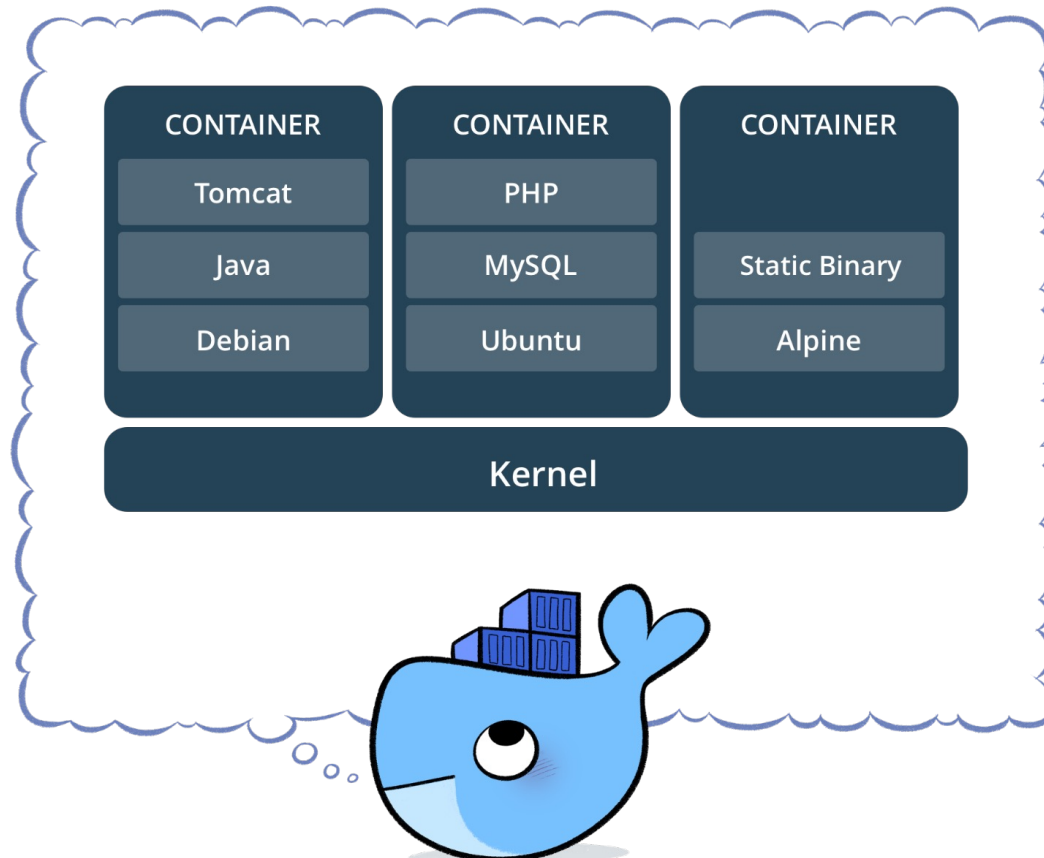

Docker



- Docker is the company driving the container movement .
- A container image is
 - A lightweight, stand-alone, executable package of a piece of software
 - It includes everything needed to run it: code, runtime, system tools, system libraries, settings.
- Available for both Linux and Windows based apps, containerized software ***will always run the same, regardless of the environment.***



Docker



Dockerfile example

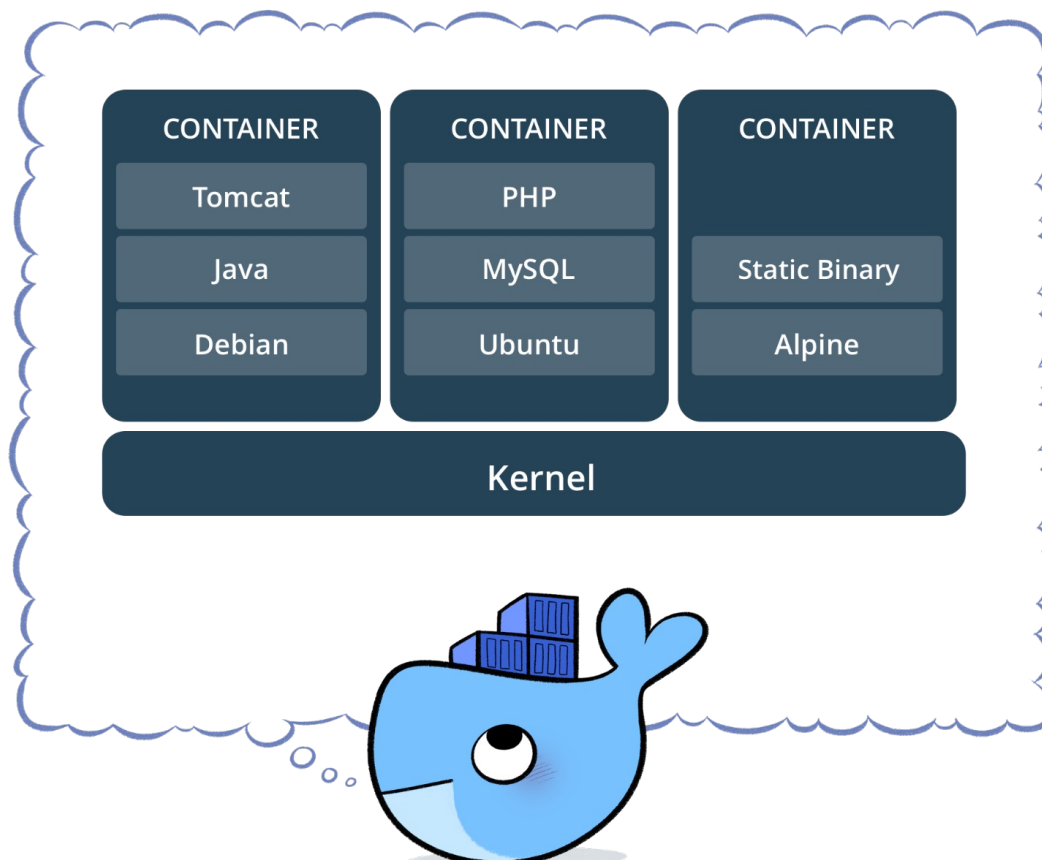
```
FROM alpine:3.4  
  
RUN apk update  
RUN apk add vim  
RUN apk add curl
```

<https://takacsmark.com/dockerfile-tutorial-by-example-dockerfile-best-practices-2018/>

Docker



docker



What about running Linux docker image in Windows?

Windows Subsystem for Linux (WSL)

- “A full Linux kernel built by Microsoft, allowing Linux distributions to run without having to manage Virtual Machines.”
- “With Docker Desktop running on WSL 2, users can leverage Linux workspaces and avoid having to maintain both Linux and Windows build scripts.”
- “In addition, WSL 2 provides improvements to file system sharing, boot time, and allows access to some cool new features for Docker Desktop users.”

<https://docs.docker.com/desktop/windows/wsl/>

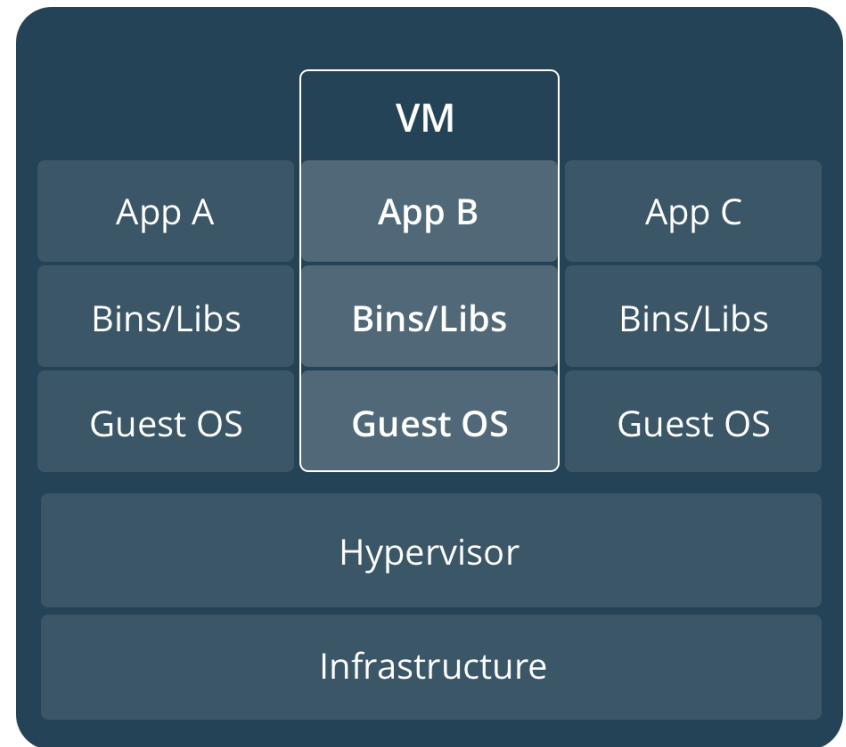
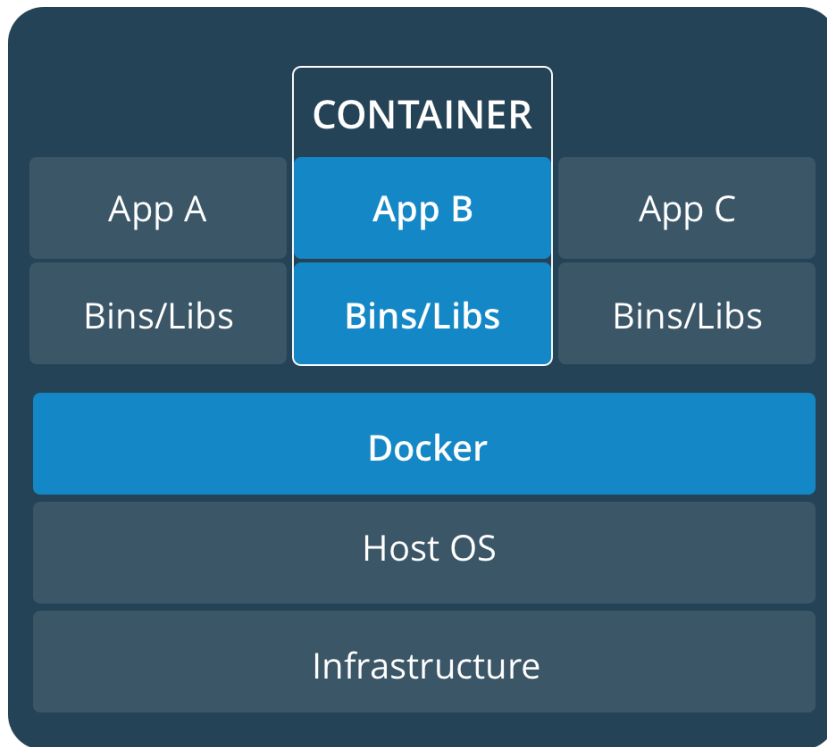
Containers vs. Virtual Machines

- Virtual machines (VMs) are an abstraction of physical hardware turning one server into many servers.
- A VM includes a full copy of an operating system, one or more apps, necessary binaries & libraries-taking up tens of GBs.
- VMs can also be ***slow to boot***.

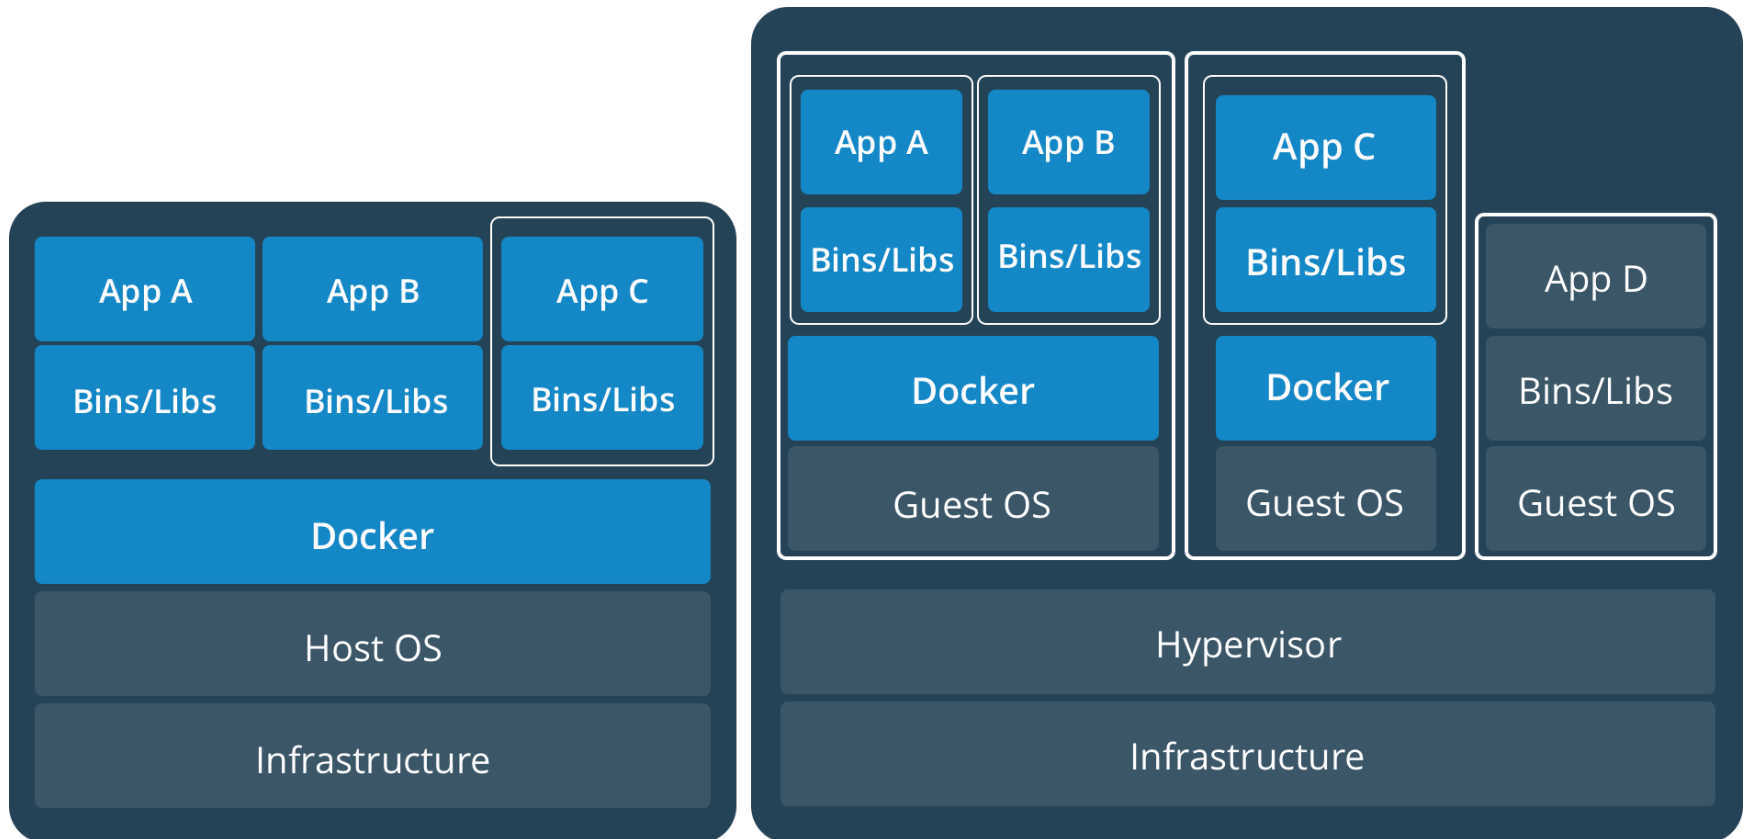
Containers vs. Virtual Machines (Cont.)

- Containers are ***an abstraction at the app layer*** that packages code and dependencies together.
- Multiple containers can run on the same machine **and share the OS kernel with other containers.**
 - each running as isolated processes in user space.
- Containers take up less space than VMs
 - Container images are typically tens of MBs in size
 - ***Start almost instantly.***

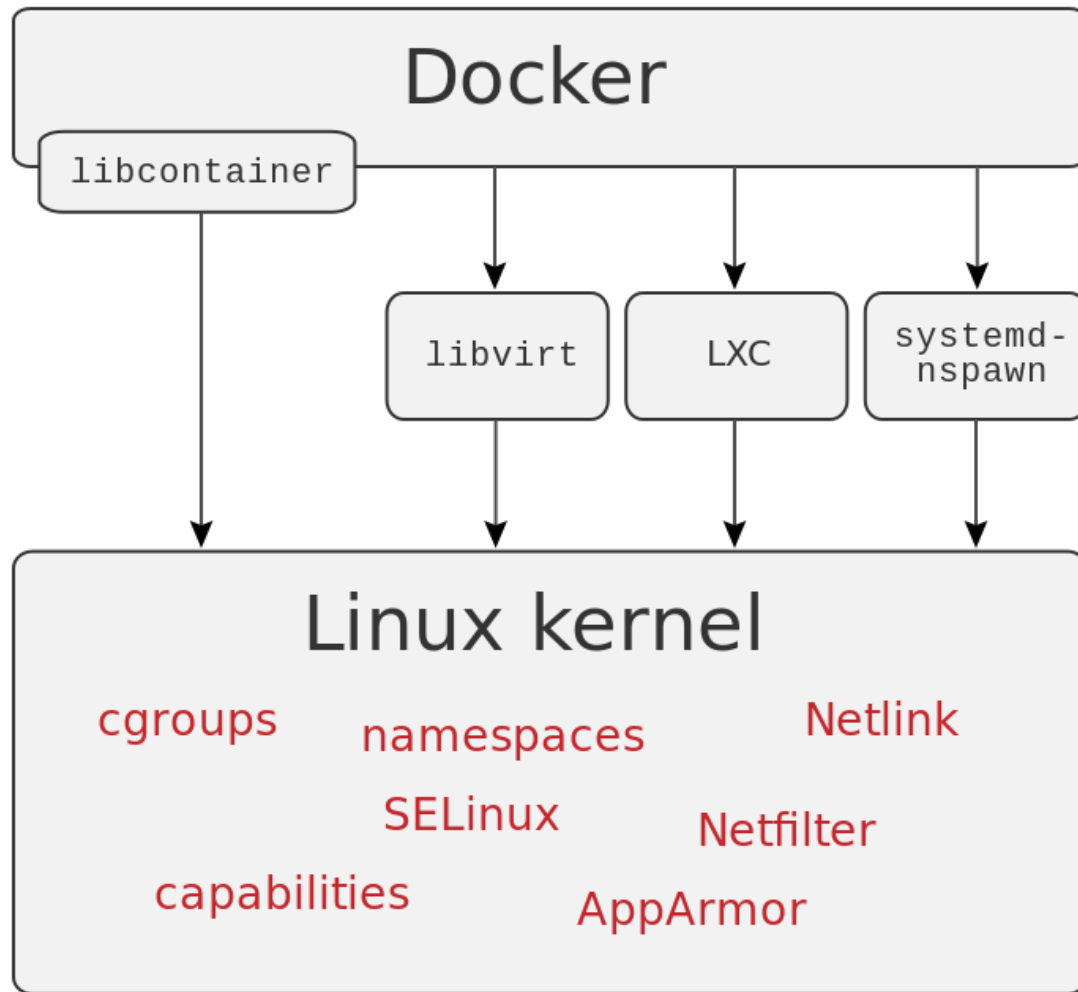
Containers vs. Virtual Machines (Cont.)



Containers vs. Virtual Machines (Cont.)




Docker Technology



Docker Technology

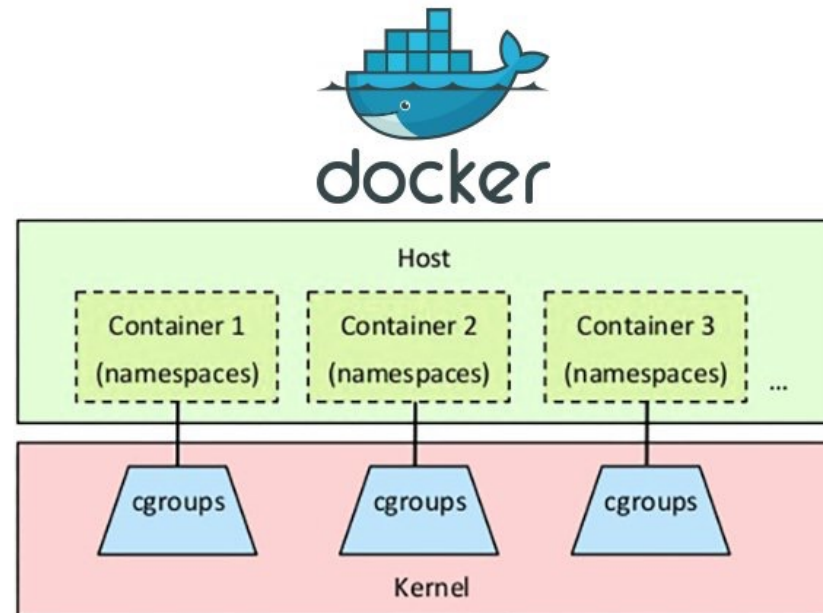
- Docker is developed ***primarily for Linux***, where it uses the resource isolation features of the Linux kernel such as:
- **cgroups and kernel namespaces,**
 - and a **union-capable file system such as OverlayFS and others.**



<https://medium.com/@knoldus/unionfs-a-file-system-of-a-container-2136cd11a779>


Docker Technology- cgroups

- **cgroups** (abbreviated from **control groups**) is a Linux kernel feature that limits, accounts for, and isolates the resource usage (CPU, memory, disk I/O, network, etc.) of a collection of processes.



Docker Technology- Namespaces

- Feature of the Linux kernel that partitions kernel resources such that one set of processes sees one set of resources while another set of processes sees a different set of resources.



<https://wvi.cz/diyC/namespaces/>

<https://www.nginx.com/blog/what-are-namespaces-cgroups-how-do-they-work/>